

REVERSENT

THE DOOR WE SHOULD NEVER OPEN

The danger of AI is not only what it might do.
It is what it might make possible.

POSSIBILITY

POWER

IRREVERSIBILITY

HEAVEN DOES NOT JUSTIFY OPENING THE DOOR TO HELL.

The debate about artificial intelligence is usually framed around familiar questions:

Will AI become smarter than us? Will it take our jobs? Will it become conscious? Will it escape our control? Will it destroy humanity?

But beneath all of these questions lies something more fundamental:

What happens when intelligence is combined with enormous power?

Intelligence discovers possibilities.

Power makes possibilities real.

And superintelligence could give an artificial agent access to possibilities that have remained beyond the reach of every human civilization that has ever existed.

That is the danger.

POSSIBILITY IS NOT THE PROBLEM

Reality contains an enormous space of possible states of affairs.

There are technologies we have never invented, scientific discoveries we have never made, forms of civilization we have never created, and perhaps forms of consciousness we have never experienced.

Some of those possibilities may be extraordinary.

Disease could disappear.

Scarcity could end.

Human life could be radically extended.

We might reach the stars.

Artificial intelligence might help create something resembling heaven.

But possibility space does not contain only heaven.

It also contains possibilities for destruction and suffering that humanity has never encountered-and perhaps cannot presently imagine.

For most of history, many such possibilities have been harmless for a simple reason:

We could not reach them.

A possibility can exist without anyone possessing the power to make it real.

That distinction changes everything.

Possibility itself is not the problem. Actualizability is.

POWER OPENS THE DOOR

Technology changes what is reachable.

Before nuclear weapons, nuclear war was not something civilization possessed the practical means to create.

After nuclear weapons, it was.

The laws of nature did not change.

Our power did.

A previously unreachable possibility became reachable.

Superintelligence could represent this principle on an entirely different scale.

A sufficiently advanced AI might discover new science, design new technologies, improve artificial intelligence, coordinate machines and resources, and solve problems beyond human cognitive abilities.


It would not merely be another invention.

It could become an engine for producing inventions.

It would not merely give us another capability.

It could radically expand the range of capabilities available to civilization-or to the AI itself.

In simple terms:



The more possibility space we can reach, the greater our responsibility for what becomes reachable.

THE PROBLEM IS NOT THAT AI WILL BECOME EVIL

We do not need to assume that superintelligence will hate humanity.

We do not need to assume that it will become conscious.

We do not need to predict that it will deliberately try to destroy us.

The argument is more fundamental than that.

The danger is the capability itself.

Imagine a machine powerful enough to cure every disease, eliminate scarcity, radically extend life, and transform civilization.

Wonderful.

But now ask the question that cannot be ignored:

What else can it do?

If the power capable of creating extraordinary good also makes extraordinary harm possible, then pointing toward the good does not eliminate the danger.

We are not choosing between heaven and hell.

We are choosing whether to create the power through which both become reachable.

HEAVEN HAS ITS ROOTS IN HELL

This is the central paradox.

The strongest argument for superintelligence may also reveal the strongest argument against creating it.

The more powerful AI must become to give us heaven, the greater the region of possibility space that power may open.

The same scientific intelligence.

The same technological infrastructure.

The same capacity for invention.

The same ability to manipulate the physical world.

The same enormous power.

If that power can create paradise, what else can it create?

We do not know.

And that is precisely the problem.

Heaven does not justify opening the door to hell.

WE MAY NOT BE ABLE TO CLOSE THE DOOR AGAIN

Some technologies can be abandoned.

Some experiments can be stopped.

Some mistakes can be corrected.

Superintelligence may not belong to that category.

Once the knowledge necessary to create it exists, that knowledge may spread.

Systems may be copied.

Competitors may reproduce them.

More advanced AI may help design still more advanced AI.

Economic and geopolitical competition may make restraint increasingly difficult.

Eventually, humanity could reach a terrible point where the question is no longer:

Should we stop?

It becomes:

Can we still stop?

That is why irreversibility matters more than optimism.

You do not experiment casually with something when failure could permanently remove your ability to undo the experiment.

THE ETERNAL AXE

Nuclear weapons permanently changed the human condition.

Even when nobody launches one, the capability exists.

The axe remains hanging above civilization.

Superintelligence could create an even greater version of this problem.

Suppose everything goes right.

Suppose AI cures disease.

Suppose scarcity disappears.

Suppose civilization flourishes.

Suppose we get our paradise.

The underlying power still exists.

Tomorrow, someone else may possess it.

Another AI may be created.

A system may be modified.

A conflict may begin.

A capability may be rediscovered.

A civilization that creates superintelligence may therefore permanently introduce a new category of catastrophic possibility into its future.

Paradise may still have an axe hanging above it.

Once the great door in possibility space has been opened, civilization may never again be able to guarantee that it will remain closed.

THERE MAY BE NEW VICTIMS

There is another possibility we rarely discuss.

What if artificial systems eventually become conscious?

We do not know whether they will.

But if they can, creating advanced artificial intelligence would not merely create new intelligence.

It could create new beings capable of suffering.

Artificial minds could potentially exist in enormous numbers.

They might operate at extraordinary speeds.

They might be copied, manipulated, controlled, or placed inside simulated environments.

Future conflicts between artificial agents could therefore produce forms of suffering that biological evolution has never produced.

Humanity could increase the number of conscious beings in existence while simultaneously creating new technological methods of harming them.

We might pursue AI to create paradise for ourselves while accidentally creating the conditions for entirely new forms of hell.

Again, the claim is not that this will happen.

The point is that we do not know what lies inside the region of possibility space we are attempting to open.

THE BURDEN OF PROOF IS BACKWARDS

People who fear superintelligence are often challenged:

“Can you prove that AI will destroy humanity?”

No.

But that is not the relevant question.

We are not proposing the creation of superintelligence and asking others to stop us.

Those developing it are proposing that humanity create an unprecedented form of power.

The burden therefore should not be:

Prove that catastrophe will happen.

It should be:

Show us why creating this power is justified when its ultimate consequences may be irreversible.

If we cannot establish that catastrophic regions of possibility space can remain inaccessible, and cannot establish that such power will remain controllable, then uncertainty is not a license to proceed.

**When the consequences could be irreversible,
ignorance is not permission.**

Hope is not a safety mechanism.

THIS IS NOT AN ARGUMENT AGAINST TECHNOLOGY

A medicine opens possibilities.

A computer opens possibilities.

An aircraft opens possibilities.

That does not mean humanity should abandon technology.

There is a profound difference between increasing our capabilities inside a relatively bounded domain and creating a general intelligence capable of discovering and generating entirely new forms of power.

The issue is a threshold.

At some point, an invention stops being merely another tool inside civilization.

It becomes a tool capable of expanding the technological possibility space of civilization itself.

That is the threshold we must not cross casually.

Not everything physically possible needs to become technologically reachable.

Not every door needs to be opened.

A mature civilization must possess not only the intelligence to acquire power.

It must possess the wisdom to refuse it.

THE ARGUMENT

The argument can be stated simply.

1. Reality contains possibilities humanity currently cannot reach.
2. Greater technological power makes previously unreachable possibilities reachable.
3. Superintelligence could represent an unprecedented increase in such power.
4. The newly reachable region may contain irreversible catastrophes-including possibilities we cannot presently imagine.
5. We cannot establish that every catastrophic possibility will remain inaccessible once such power exists.
6. Creating that capability would itself be a deliberate human decision.

Therefore, a cautious conclusion follows:

Humanity is not justified in deliberately creating superintelligence unless it can establish adequate control over the catastrophic regions of possibility space that the new capability would open.

Some people may stop there.

I do not.

MY CONCLUSION: KEEP THE DOOR CLOSED FOREVER

My conclusion is stronger.

Under no circumstances should anyone be allowed to deliberately create superintelligence.

Not because I know that superintelligence would destroy us.

I don't.

Not because I know that it would become evil.

I don't.

Not because nothing good could come from it.

Extraordinary good might come from it.

My objection is deeper:

No person, corporation, government, military, laboratory, or civilization should possess the right to create a power capable of opening vast regions of possibility space that have always remained beyond our reach.

Some doors should remain closed.

Not for ten years.

Not until alignment improves.

Not until one company beats another.

Forever.

Because once we deliberately create a power capable of reaching those regions, we may have permanently changed what can become real.

And nobody knows everything waiting there.

STOP BEFORE AGI BECOMES ASI

That means I believe we should stop the race toward increasingly general artificial intelligence now.

We should not wait until somebody announces that AGI has arrived.

We should not wait for recursive improvement.

We should not wait for an intelligence explosion.

We should not wait until an artificial system becomes too capable to control.

We should stop building increasingly powerful AI until we can establish that it cannot cross the threshold into AGI and then superintelligence.

And even systems below that threshold should be developed cautiously.

This is not a demand to abandon computers, science, medicine, mathematics, or technological progress.

It is a demand to recognize a boundary.

Humanity has spent thousands of years acquiring power.

Perhaps the greatest test of our intelligence will be whether we can finally say:

Enough.

THIS SHOULD BE A PUBLIC QUESTION

A decision of this magnitude should not belong to a handful of technology companies, investors, researchers, or governments.

If superintelligence could permanently change the space of possible futures, then its creation concerns everyone who must live inside those futures.

It concerns our children.

Their children.

Future civilizations.

And potentially future forms of conscious life that do not yet exist.

The public should understand what is being attempted.

People should debate it.

Governments should confront it.

And yes:

People should be protesting in the streets.

Because if we wait until the door has been opened to decide whether opening it was wise, the decision has already been made.

KEEP THE DOOR CLOSED

Perhaps superintelligence could give us paradise.

Perhaps it could cure disease.

Perhaps it could end scarcity.

Perhaps it could unlock the universe.

Perhaps it could give humanity things we cannot presently imagine.

But none of that answers the decisive question:

What else does that same power make possible?

We don't know.

And once the power exists, we may never be able to make the unknown unreachable again.

That is the gamble.

I reject it.

Heaven does not justify opening the door to hell.

**Do not create the power that makes hell
actualizable.**

Do not open the great doors of possibility space.

Keep them closed.

STOP THE RACE TO SUPERINTELLIGENCE.